gz-unitree: Reinforcement learning en robotique avec validation par moteurs de physique multiples pour le robot *H1v2* d'Unitree

Gwenn Le Bihan <gwenn.lebihan@etu.inp-n7.fr>

6 Novembre 2025

Reinforcement Learning

Et son application à la robotique

Bases du RL

Agent Environnement Score

Bases du RL

$$\text{Agent} \xrightarrow{\text{Action}} \text{Environnement} \qquad \text{Score}$$

Bases du RL

Bases du RL

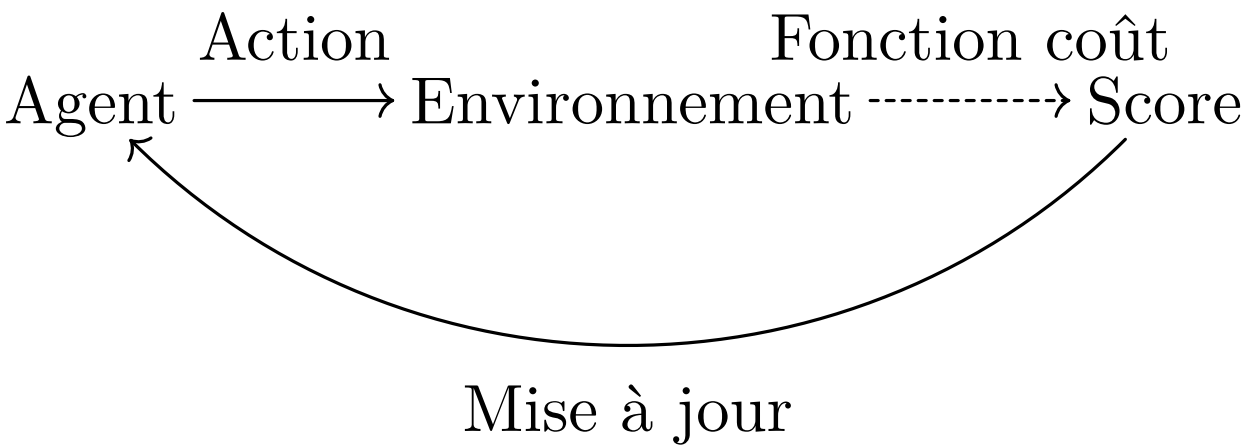
Action
Fonction coût
Agent
Environnement
Score
Mise à jour

RL en robotique

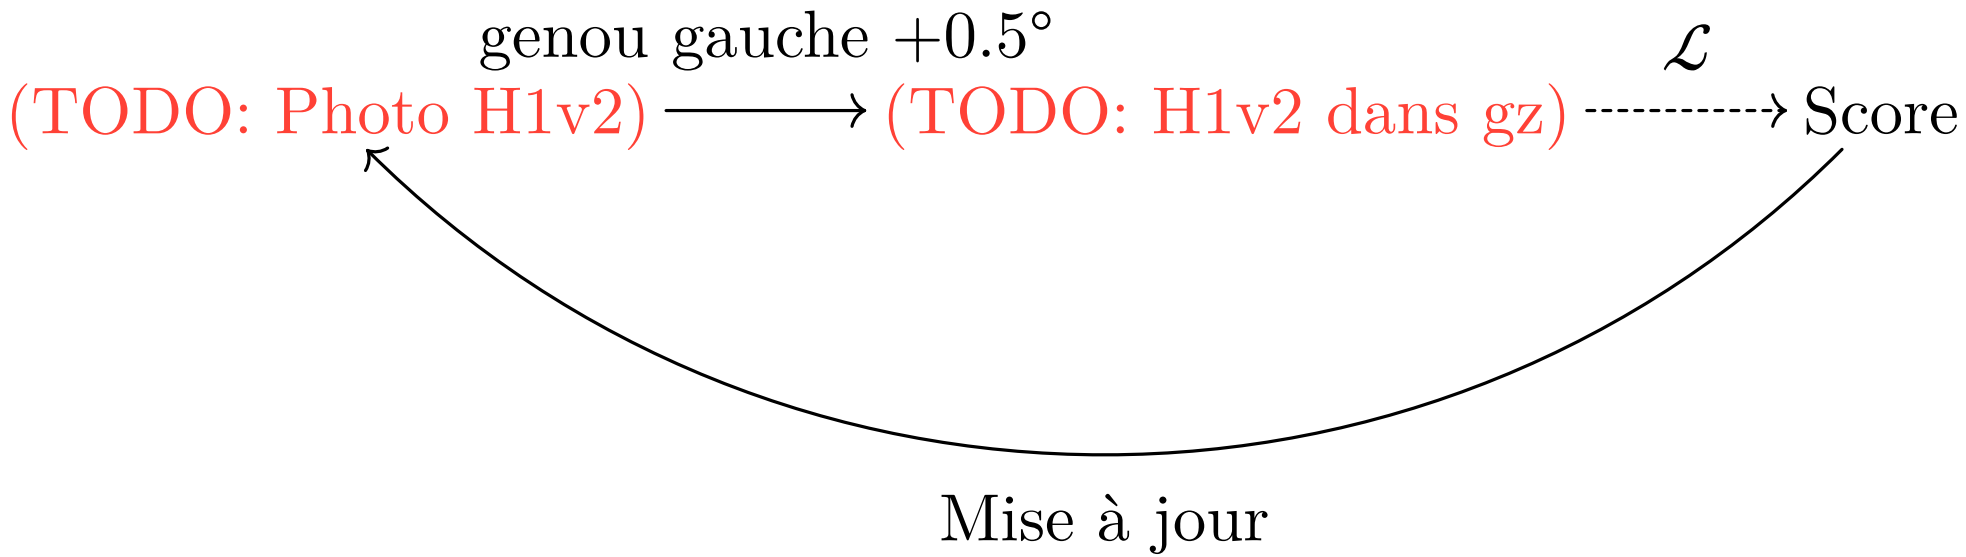

C'est quoi $\mathcal{L}$?

C'est très simple:

$$\mathcal{L}_r(\pi', \pi) := \mathop{\mathbb{E}}_{(s_t, a_t)_{t\in\mathbb{N}} \in \mathcal{C}} \sum_{t=0}^{\infty} \frac{Q_\pi(s_t, a_t)}{Q_{\pi'}(s_t, a_t)} A_{\pi,r}(s_t, a_t)$$

Comparaison des politiques

En Reinforcement Learning

Comparaison des politiques

$$s_t$$

Comparaison des politiques

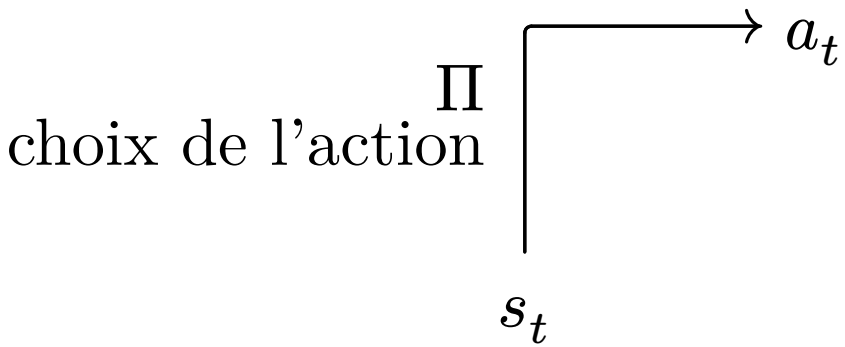

Comparaison des politiques

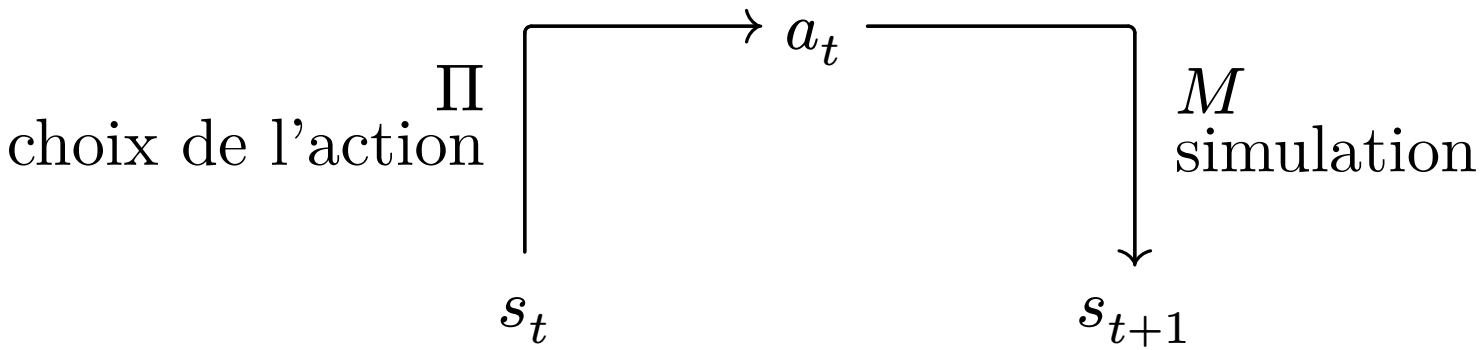

Comparaison des politiques

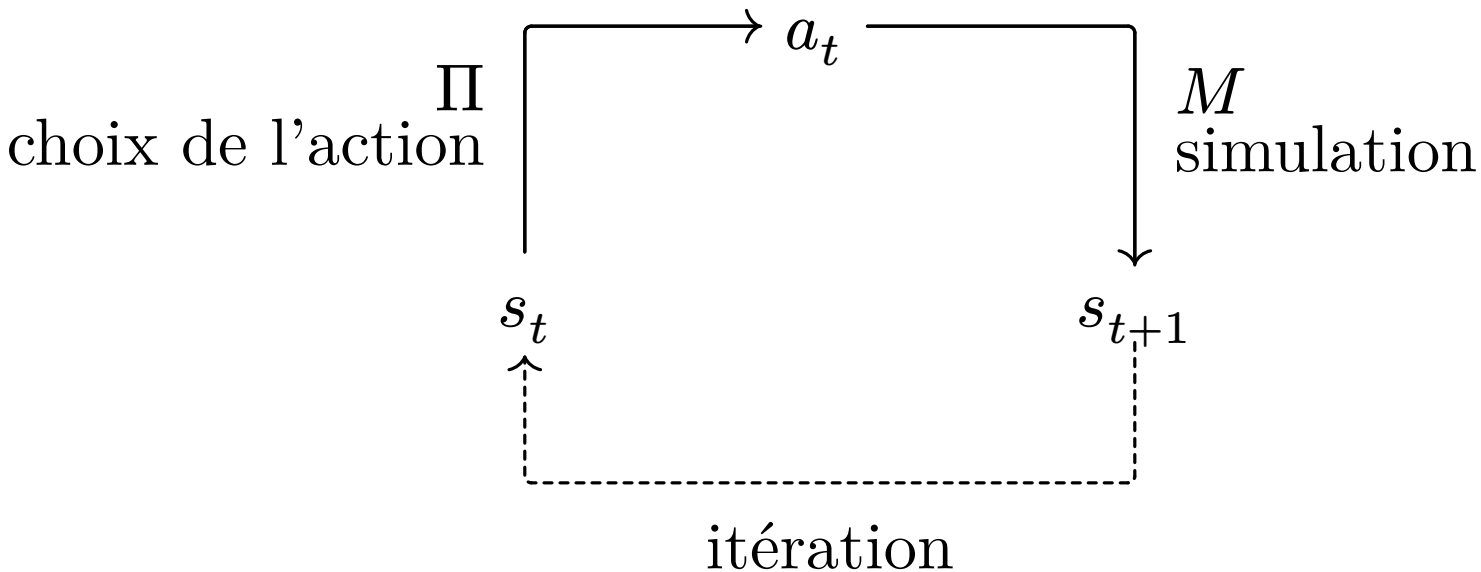

Comparaison des politiques

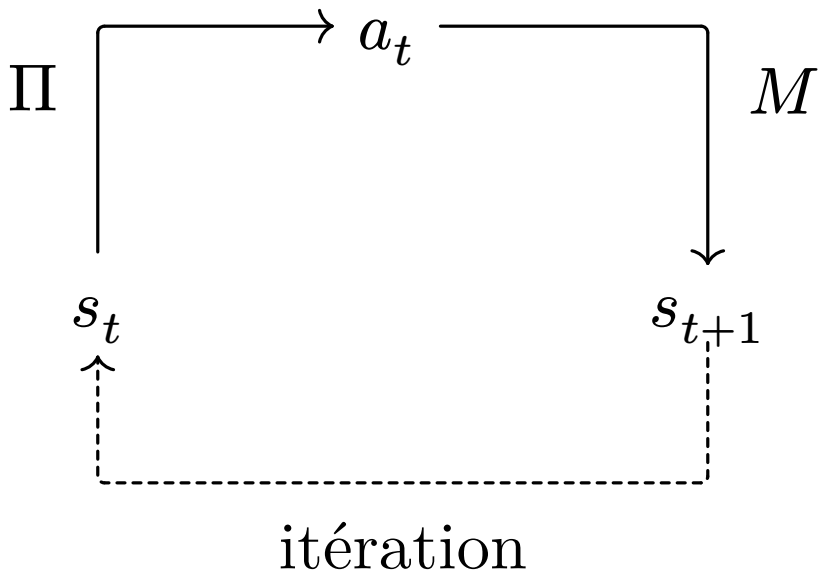

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} \cdots$$

Comparaison des politiques

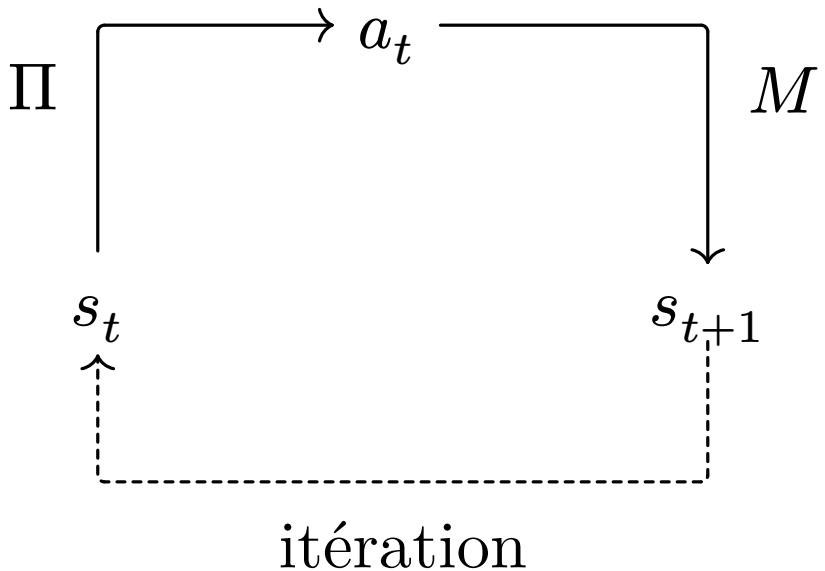

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} \cdots$$

$$((s_0, a_0), (s_1, a_1), (s_2, a_2), ...)$$

Comparaison des politiques

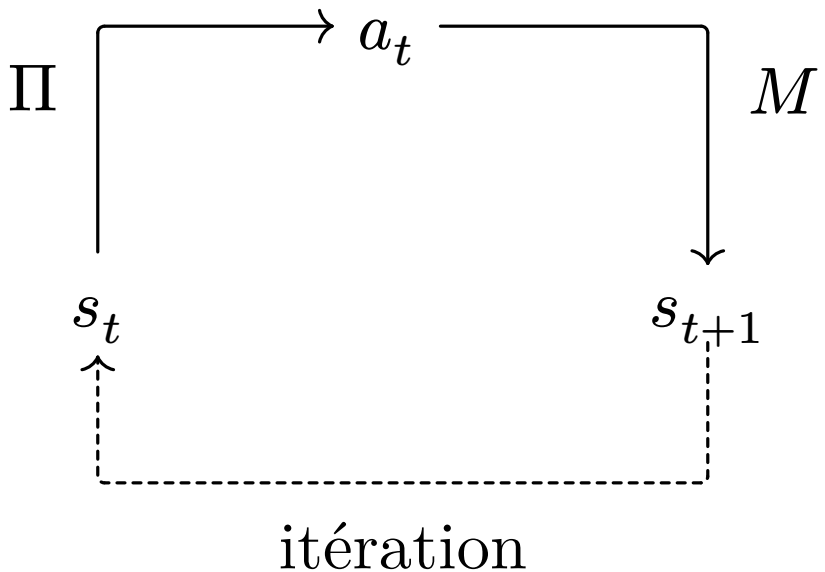

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} \cdots$$

$$((s_0, a_0), (s_1, a_1), (s_2, a_2), ...) \in \mathcal{C}$$

Comparaison des politiques

$A :=$ actions possibles

$S :=$ états possibles

$$\mathcal{C} := \left\{ \left\{ \begin{array}{ll} & c_0 = (s_0, a_0) \\ \forall t \in \mathbb{N} & c_{t+1} = (M(c_t), a_t) \end{array} \right| (s_0, a) \in S \times A^{\mathbb{N}} \right\}$$

Comparaison des politiques: Avantage A

À quel point est-il mieux de choisir a_t plutôt qu'une autre action?

Comparaison des politiques: Avantage A

$$\cdots \xrightarrow{a_{t-2}} s_{t-1} \xrightarrow{a_{t-1}}$$

Comparaison des politiques: Avantage A

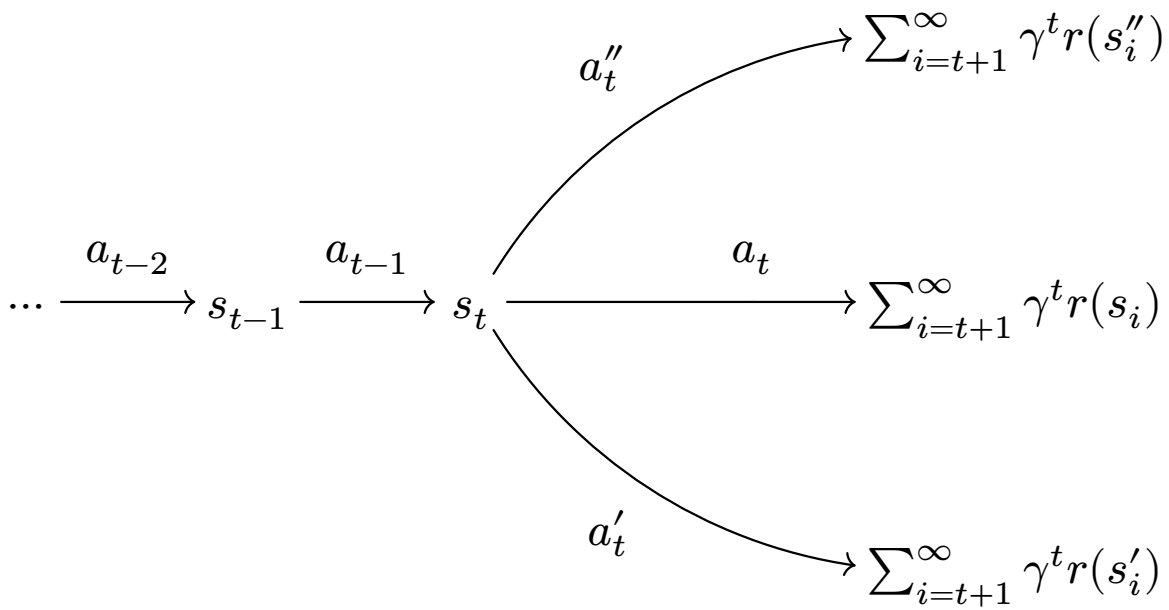

Comparaison des politiques: Avantage A

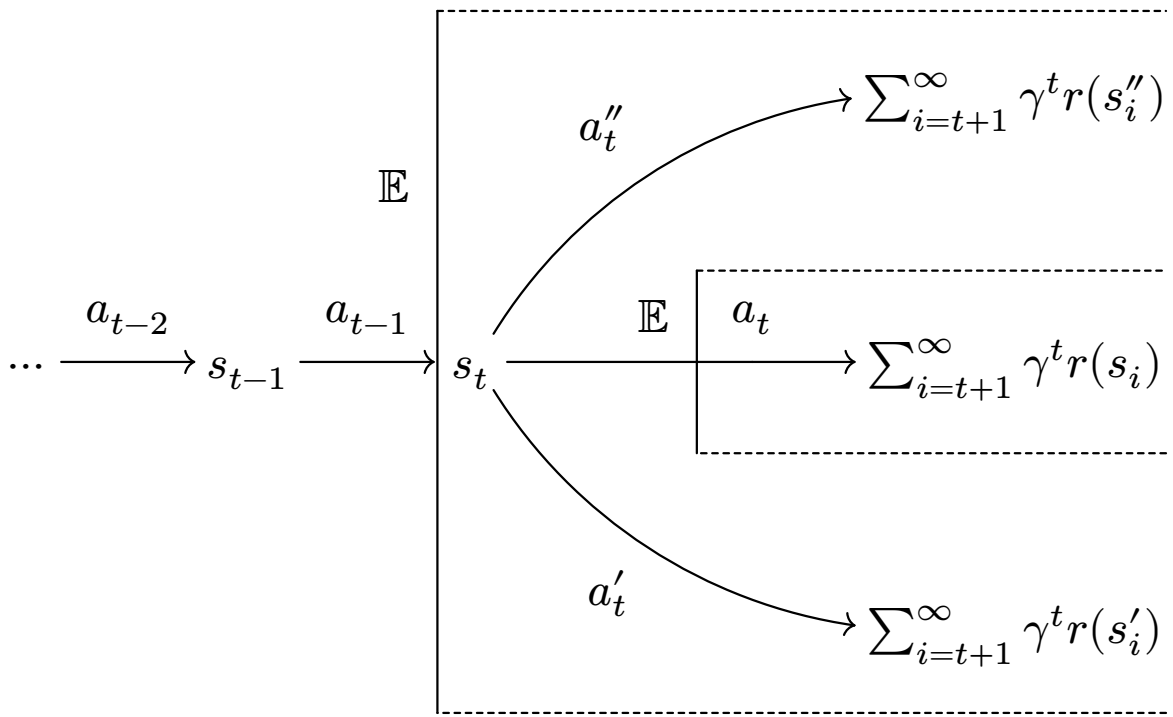

Comparaison des politiques: Avantage A

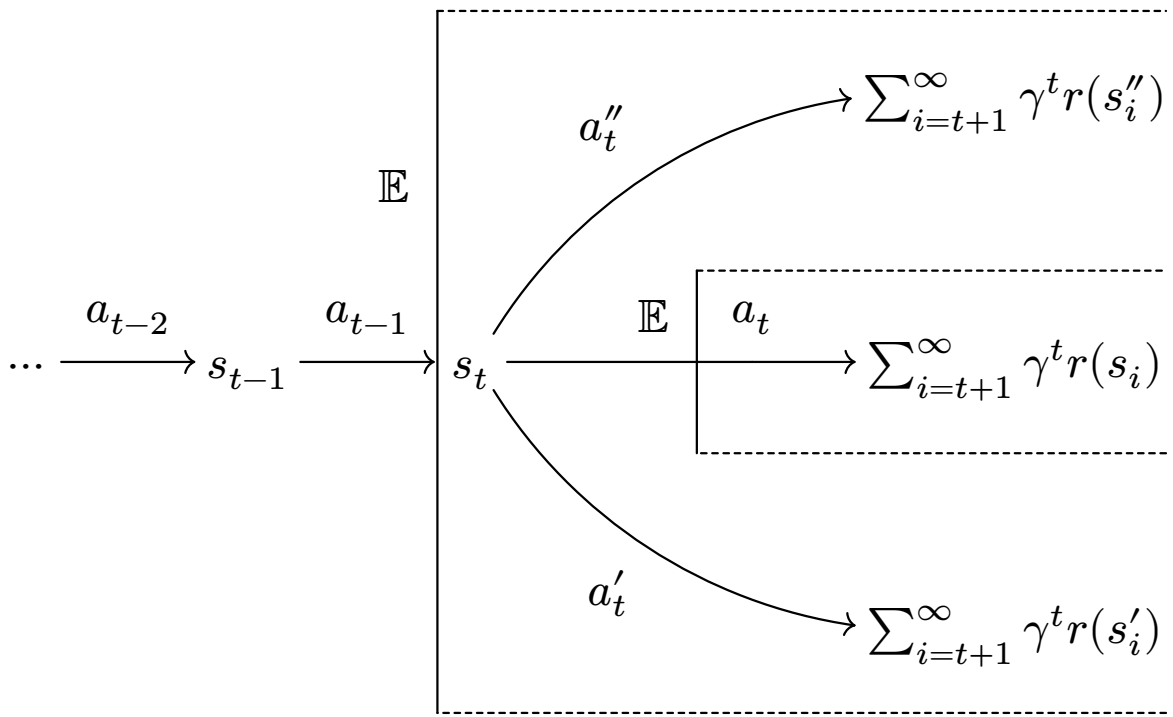

$$A_{\pi,r}(s,a) := \mathbb{E}(\text{avec } a_t) - \mathbb{E}(\text{à } t-1)$$

C'est quoi $\mathcal{L}$?

$$\mathcal{L}_r(\pi', \pi) :=$$

C'est quoi $\mathcal{L}$?

$$\mathcal{L}_r(\pi', \pi) := \mathop{\mathbb{E}}_{(s_t, a_t)_{t \in \mathbb{N}} \in \mathcal{C}}$$

C'est quoi $\mathcal{L}$?

$$\mathcal{L}_r(\pi', \pi) := \mathop{\mathbb{E}}_{(s_t, a_t)_{t \in \mathbb{N}} \in \mathcal{C}} \sum_{t=0}^{\infty}$$

C'est quoi $\mathcal{L}$?

$$\mathcal{L}_r(\pi', \pi) := \mathop{\mathbb{E}}_{(s_t, a_t)_{t\in\mathbb{N}} \in \mathcal{C}} \sum_{t=0}^{\infty} \frac{Q_\pi(s_t, a_t)}{Q_{\pi'}(s_t, a_t)} A_{\pi,r}(s_t, a_t)$$

Mise à jour de Π

$$\Pi' = \begin{cases} \arg\max_{\pi} \mathcal{L}_r(\pi, \Pi) \\ \text{s.c.} \quad \text{distance}(\Pi', \Pi) < \delta \end{cases}$$

Mise à jour de Π: distance entre politiques

$$\text{distance}(\Pi', \Pi) := \max_{s \in S} D_{\text{KL}}(Q_{\Pi'}(s, \cdot) \parallel Q_{\Pi}(s, \cdot))$$

$$D_{\text{KL}}(P \parallel P') := \sum_{x \in \mathcal{X}} P(x) \log \frac{P(x)}{P'(x)}$$

Le *SDK*[1] d'Unitree

[1]Software Development Kit

DDS

overlayed-img[

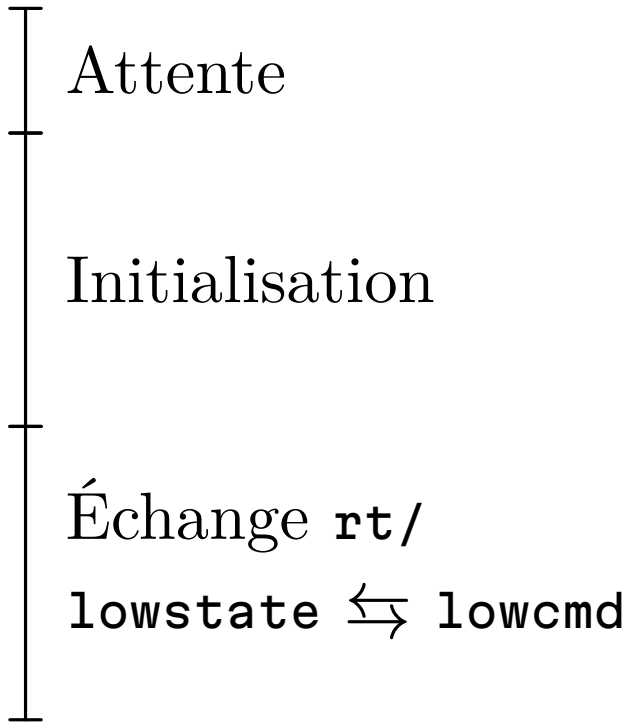
Attente
Initialisation
Échange rt/
lowstate ⇆ lowcmd

],

Le SDK d'Unitree

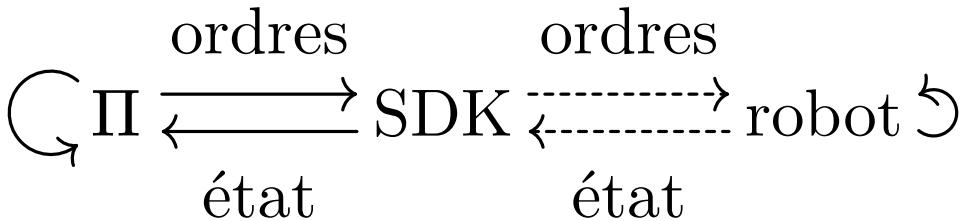

⇢ Message DDS

Le SDK d'Unitree

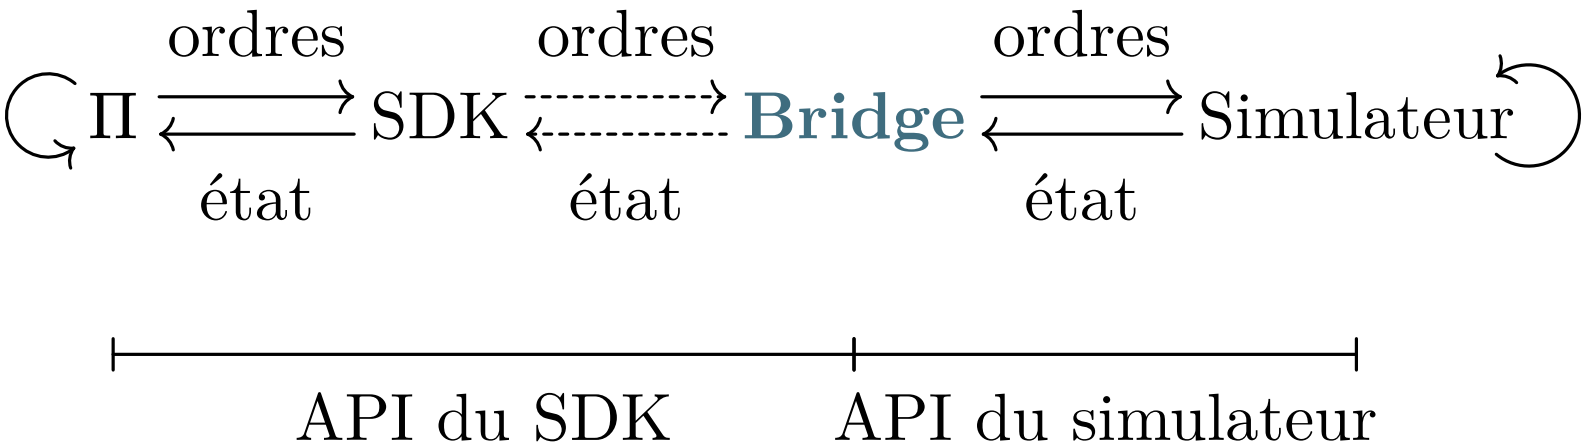

unitree_mujoco

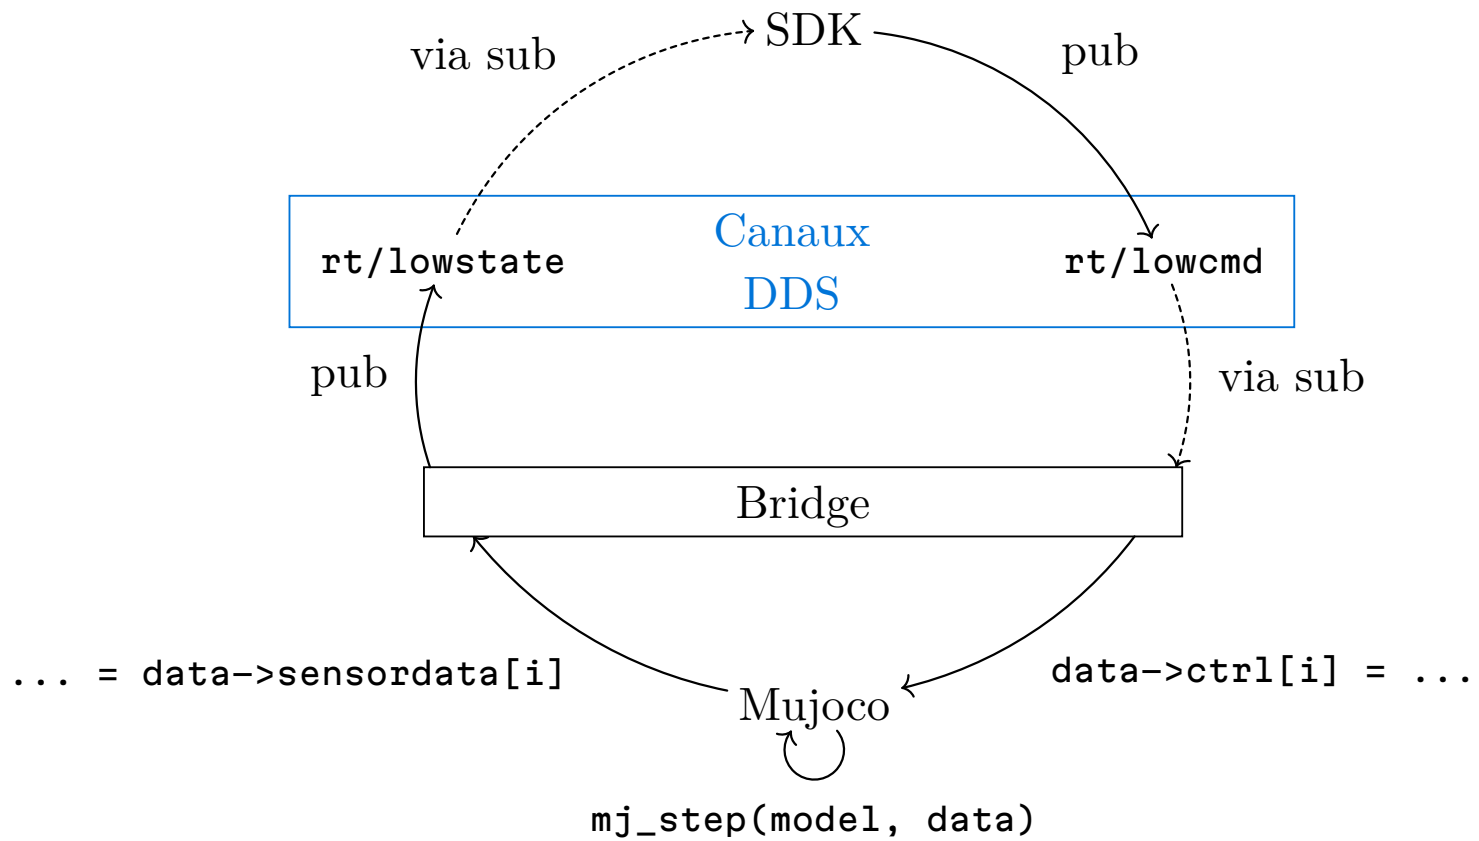

Développement de *gz-unitree*

Un bridge pour Gazebo